



МИНОБРНАУКИ РОССИИ
Федеральное государственное бюджетное образовательное учреждение
высшего образования
«МИРЭА – Российский технологический университет»

Институт кибербезопасности и цифровых технологий
Кафедра КБ-4 «Интеллектуальные системы информационной безопасности»

Отчёт по лабораторной работе № 1

По дисциплине

«Анализ защищенности систем искусственного интеллекта»

Выполнил:

БМО–02–22

Шмарковский М. Б.

Проверил:

Спирин А. А.

Москва, 2024

Каждый фрагмент кода начинается с краткого комментария, описывающего процесс, в целях которого этот фрагмент и написан.

Вывод для каждого фрагмента кода для наглядности будет написан на сером фоне.

```
# Сначала скопируем проект по ссылке в локальную среду Google Colab  
!git clone https://github.com/ewatson2/EEL6812_DeepFool_Project.git
```

```
Cloning into 'EEL6812_DeepFool_Project'...  
remote: Enumerating objects: 96, done.  
remote: Counting objects: 100% (3/3), done.  
remote: Compressing objects: 100% (2/2), done.  
remote: Total 96 (delta 2), reused 1 (delta 1), pack-reused 93  
Receiving objects: 100% (96/96), 33.99 MiB | 11.54 MiB/s, done.  
Resolving deltas: 100% (27/27), done.
```

```
# Меняем директорию исполнения на папку проекта  
%cd /content/EEL6812_DeepFool_Project
```

```
/content/EEL6812_DeepFool_Project
```

```
# Импортируем необходимые библиотеки  
import numpy as np  
import json, torch  
from torch.utils.data import DataLoader, random_split  
from torchvision import datasets, models  
from torchvision.transforms import transforms
```

```
# Импортируем вспомогательные библиотеки из файлов проекта  
from models.project_models import FC_500_150, LeNet_CIFAR, LeNet_MNIST,  
Net  
from utils.project_utils import get_clip_bounds, evaluate_attack,  
display_attack
```

```
# Установим случайное рандомное значение в виде переменной  
rand_seed={"Порядковый номер ученика группы в Гугл-таблице"}, укажем  
значение для np.random.seed и torch.manual_seed  
rand_seed = 22  
np.random.seed(rand_seed)  
torch.manual_seed(rand_seed)
```

```
<torch._C.Generator at 0x7ce54824a550>
```

```
# Используем в качестве устройства видеокарту  
use_cuda = torch.cuda.is_available()  
device = torch.device('cuda' if use_cuda else 'cpu')
```

```
# Загрузим датасет MNIST с параметрами mnist_mean = 0.5, mnist_std = 0.5,  
mnist_dim = 28  
mnist_mean = 0.5  
mnist_std = 0.5  
mnist_dim = 28
```

```
mnist_min, mnist_max = get_clip_bounds(mnist_mean,  
                                         mnist_std,
```

```

mnist_dim)

mnist_min = mnist_min.to(device)
mnist_max = mnist_max.to(device)

mnist_tf = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize(
        mean=mnist_mean,
        std=mnist_std)])

mnist_tf_train = transforms.Compose([
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize(
        mean=mnist_mean,
        std=mnist_std)])

mnist_tf_inv = transforms.Compose([
    transforms.Normalize(
        mean=0.0,
        std=np.divide(1.0, mnist_std)),
    transforms.Normalize(
        mean=np.multiply(-1.0, mnist_std),
        std=1.0)])

mnist_temp = datasets.MNIST(root='datasets/mnist', train=True,
                             download=True, transform=mnist_tf_train)
mnist_train, mnist_val = random_split(mnist_temp, [50000, 10000])

mnist_test = datasets.MNIST(root='datasets/mnist', train=False,
                              download=True, transform=mnist_tf)

cifar_classes = ['airplane', 'automobile', 'bird', 'cat', 'deer',
                  'dog', 'frog', 'horse', 'ship', 'truck']

```

```

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to da
taset/mnist/MNIST/raw/train-images-idx3-ubyte.gz
100%|██████████| 9912422/9912422 [00:00<00:00, 122629555.19it/s]
Extracting datasets/mnist/MNIST/raw/train-images-idx3-ubyte.gz to datasets/mn
ist/MNIST/raw

```

```

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to da
taset/mnist/MNIST/raw/train-labels-idx1-ubyte.gz
100%|██████████| 28881/28881 [00:00<00:00, 39677593.78it/s]
Extracting datasets/mnist/MNIST/raw/train-labels-idx1-ubyte.gz to datasets/mn
ist/MNIST/raw

```

```

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to da
taset/mnist/MNIST/raw/t10k-images-idx3-ubyte.gz
100%|██████████| 1648877/1648877 [00:00<00:00, 47687909.56it/s]
Extracting datasets/mnist/MNIST/raw/t10k-images-idx3-ubyte.gz to datasets/mni
st/MNIST/raw
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to da
taset/mnist/MNIST/raw/t10k-labels-idx1-ubyte.gz
100%|██████████| 4542/4542 [00:00<00:00, 19012503.76it/s]

```

```
Extracting datasets/mnist/MNIST/raw/t10k-labels-idx1-ubyte.gz to datasets/mnist/MNIST/raw
```

```
# Загрузим датасет CIFAR-10 с параметрами cifar_mean = [0.491, 0.482, 0.447] cifar_std = [0.202, 0.199, 0.201] cifar_dim = 32
cifar_mean = [0.491, 0.482, 0.447]
cifar_std = [0.202, 0.199, 0.201]
cifar_dim = 32
```

```
cifar_min, cifar_max = get_clip_bounds(cifar_mean,
                                       cifar_std,
                                       cifar_dim)
```

```
cifar_min = cifar_min.to(device)
cifar_max = cifar_max.to(device)
```

```
cifar_tf = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize(
        mean=cifar_mean,
        std=cifar_std)])
```

```
cifar_tf_train = transforms.Compose([
    transforms.RandomCrop(
        size=cifar_dim,
        padding=4),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize(
        mean=cifar_mean,
        std=cifar_std)])
```

```
cifar_tf_inv = transforms.Compose([
    transforms.Normalize(
        mean=[0.0, 0.0, 0.0],
        std=np.divide(1.0, cifar_std)),
    transforms.Normalize(
        mean=np.multiply(-1.0, cifar_mean),
        std=[1.0, 1.0, 1.0])])
```

```
cifar_temp = datasets.CIFAR10(root='datasets/cifar-10', train=True,
                              download=True, transform=cifar_tf_train)
cifar_train, cifar_val = random_split(cifar_temp, [40000, 10000])
```

```
cifar_test = datasets.CIFAR10(root='datasets/cifar-10', train=False,
                              download=True, transform=cifar_tf)
```

```
Downloading https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz to datasets/cifar-10/cifar-10-python.tar.gz
```

```
100%|██████████| 170498071/170498071 [00:03<00:00, 43315076.71it/s]
```

```
Extracting datasets/cifar-10/cifar-10-python.tar.gz to datasets/cifar-10
Files already downloaded and verified
```

```
# Выполним настройку и загрузку DataLoader batch_size = 64 workers = 4
batch_size = 64
workers = 4
```

```
mnist_loader_train = DataLoader(mnist_train, batch_size=batch_size,
                                shuffle=True, num_workers=workers)
mnist_loader_val = DataLoader(mnist_val, batch_size=batch_size,
                              shuffle=False, num_workers=workers)
mnist_loader_test = DataLoader(mnist_test, batch_size=batch_size,
```

```

        shuffle=False, num_workers=workers)

cifar_loader_train = DataLoader(cifar_train, batch_size=batch_size,
                                shuffle=True, num_workers=workers)
cifar_loader_val = DataLoader(cifar_val, batch_size=batch_size,
                              shuffle=False, num_workers=workers)
cifar_loader_test = DataLoader(cifar_test, batch_size=batch_size,
                               shuffle=False, num_workers=workers)

/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:557: UserWarning: This DataLoader will create 4 worker processes in total. Our suggested max number of worker in current system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.
  warnings.warn(_create_warning_msg(

import os
train_model = True

epochs = 50
epochs_nin = 100

lr = 0.004
lr_nin = 0.01
lr_scale = 0.5

momentum = 0.9

print_step = 5

deep_batch_size = 64
deep_num_classes = 10
deep_overshoot = 0.02
deep_max_iters = 50

deep_args = [deep_batch_size, deep_num_classes,
              deep_overshoot, deep_max_iters]

if not os.path.isdir('weights/deepfool'):
    os.makedirs('weights/deepfool', exist_ok=True)

if not os.path.isdir('weights/fgsm'):
    os.makedirs('weights/fgsm', exist_ok=True)

# Загрузим и оценим стойкость модели Network-In-Network Model к FGSM и
# DeepFool атакам на основе датасета CIFAR-10
fgsm_eps = 0.6
model = LeNet_MNIST().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_lenet.pth', map_location=torch.device('cpu')))
evaluate_attack('mnist_lenet_fgsm.csv',
                'results', device, mnist_loader_test,
                mnist_min, mnist_max, fgsm_eps, is_fgsm=True)

print('')

evaluate_attack('mnist_lenet_deepfool.csv', 'results', device, model,
mnist_loader_test, mnist_min, mnist_max, deep_args, is_fgsm=False)
if device.type == 'cuda': torch.cuda.empty_cache()

```

```
FGSM Test Error : 87.89%
FGSM Robustness : 4.58e-01
FGSM Time (All Images) : 0.29 s
FGSM Time (Per Image) : 28.86 us
```

```
DeepFool Test Error : 98.74%
DeepFool Robustness : 9.64e-02
DeepFool Time (All Images) : 193.32 s
DeepFool Time (Per Image) : 19.33 ms
```

```
# Загрузим и оценим стойкость модели LeNet к FGSM и DeepFool атакам на основе датасета CIFAR-10
```

```
fgsm_eps = 0.2
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth',
map_location=torch.device('cpu'))))

evaluate_attack('mnist_fc_fgsm.csv', 'results', device, model,
mnist_loader_test, mnist_min, mnist_max, fgsm_eps,
is_fgsm=True)
print('')
```

```
evaluate_attack('mnist_fc_deepfool.csv', 'results', device, model,
mnist_loader_test, mnist_min, mnist_max, deep_args, is_fgsm=False)
if device.type == 'cuda': torch.cuda.empty_cache()
```

```
FGSM Test Error : 87.08%
FGSM Robustness : 1.56e-01
FGSM Time (All Images) : 0.15 s
FGSM Time (Per Image) : 14.99 us
```

```
DeepFool Test Error : 97.92%
DeepFool Robustness : 6.78e-02
DeepFool Time (All Images) : 141.81 s
DeepFool Time (Per Image) : 14.18 ms
```

```
# Загрузим и оценим стойкость модели LeNet к FGSM и DeepFool атакам на основе датасета CIFAR-10
```

```
fgsm_eps = 0.1
model = LeNet_CIFAR().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_lenet.pth',
map_location=torch.device('cpu'))))

evaluate_attack('cifar_lenet_fgsm.csv', 'results', device, model,
cifar_loader_test, cifar_min, cifar_max, fgsm_eps, is_fgsm=True)
print('')
evaluate_attack('cifar_lenet_deepfool.csv', 'results', device, model,
cifar_loader_test, cifar_min, cifar_max, deep_args, is_fgsm=False)

if device.type == 'cuda': torch.cuda.empty_cache()
```

```
FGSM Test Error : 91.71%
FGSM Robustness : 8.90e-02
FGSM Time (All Images) : 0.40 s
FGSM Time (Per Image) : 40.08 us
```

```
DeepFool Test Error : 87.81%
DeepFool Robustness : 1.78e-02
DeepFool Time (All Images) : 73.27 s
```

DeepFool Time (Per Image) : 7.33 ms

Выполним оценку атакующих примеров для сетей

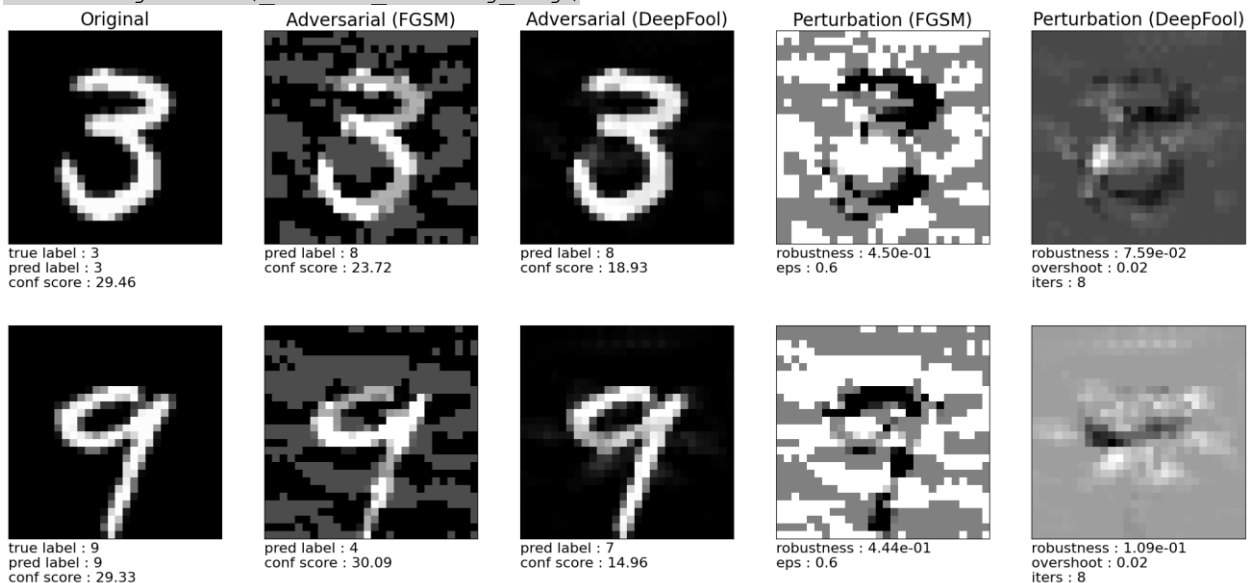
```
# LeNet на датасете MNIST
```

```
fgsm_eps = 0.6
model = LeNet_MNIST().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_lenet.pth',
map_location=device))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min,
mnist_max, fgsm_eps, deep_args,
                has_labels=False, l2_norm=True, pert_scale=1.0, fig_rows=2,
fig_width=25, fig_height=11)
```

```
if device.type == 'cuda': torch.cuda.empty_cache()
```

```
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:557: UserWarning: This DataLoader will create 4 worker processes in total. Our suggested max number of worker in current system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.
```

```
warnings.warn(_create_warning_msg(
```

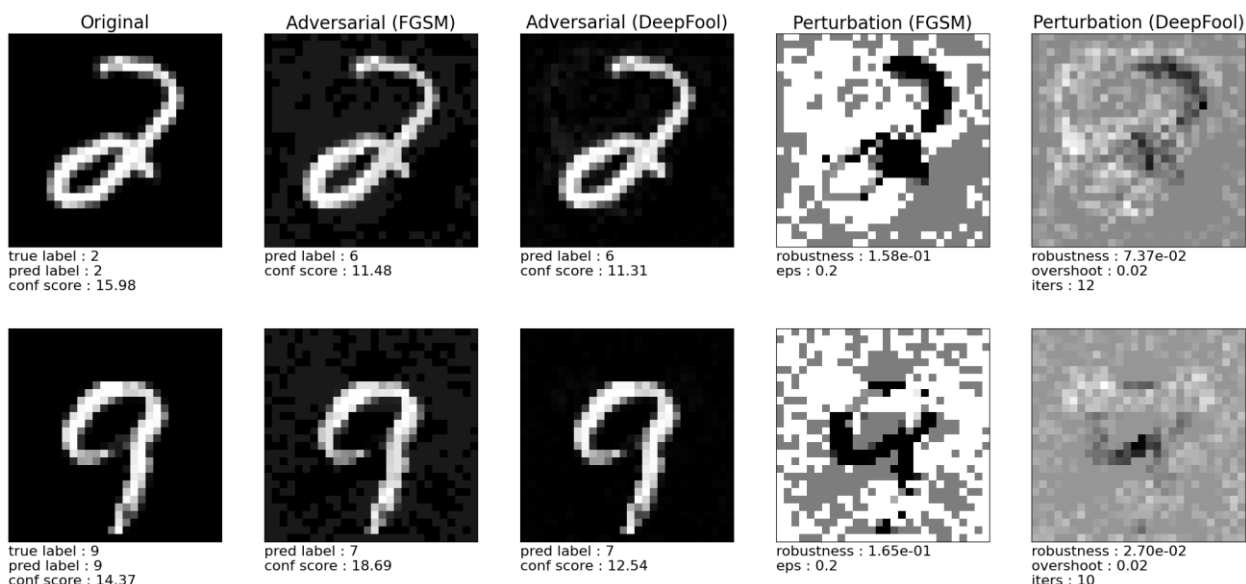


```
# FCNet на датасете MNIST
```

```
fgsm_eps = 0.2
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
```

```
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min,
mnist_max, fgsm_eps, deep_args, has_labels=False, l2_norm=True,
pert_scale=1.0, fig_rows=2, fig_width=25, fig_height=11)
```

```
if device.type == 'cuda': torch.cuda.empty_cache()
```



In [19]:

```
# Network-in-Network на датасете CIFAR
```

```
fgsm_eps = 0.2
```

```
model = Net().to(device)
```


```
model.load_state_dict(torch.load('weights/clean/cifar_nin.pth'))
```

```
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min,
cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=True,
pert_scale=1.0, fig_rows=2, fig_width=25, fig_height=11,
label_map=cifar_classes)
```

```
if device.type == 'cuda': torch.cuda.empty_cache()
```

```
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:557: UserWarning: This DataLoader will create 4 worker processes in total. Our suggested max number of worker in current system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.
```

```
warnings.warn(_create_warning_msg(
```



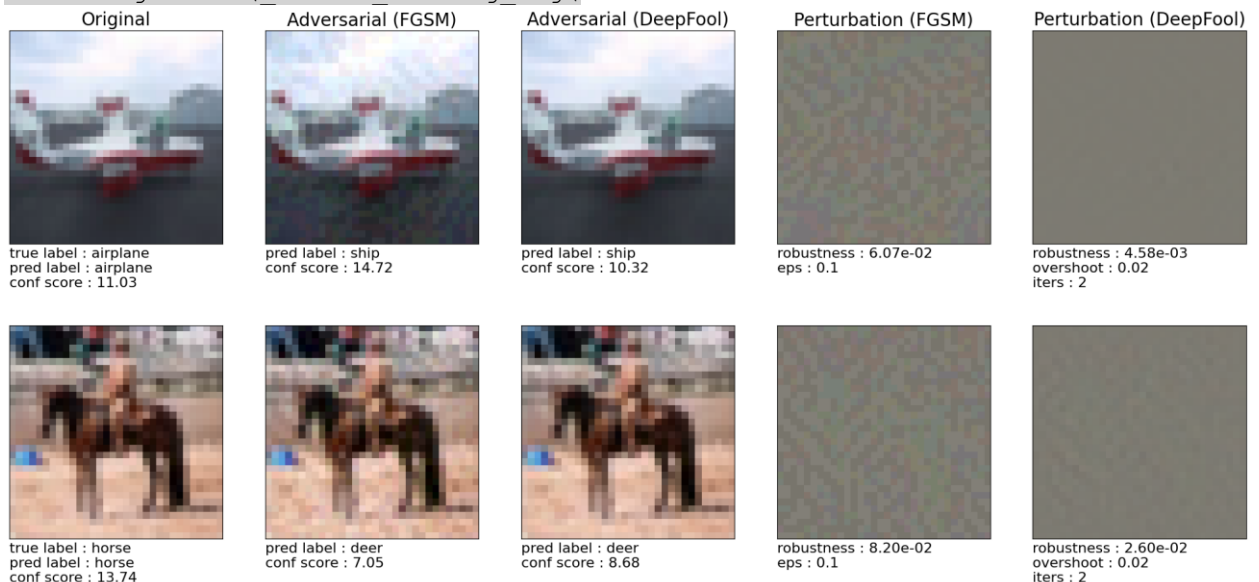

```
# LeNet на датасете CIFAR
fgsm_eps = 0.1
model = LeNet_CIFAR().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_lenet.pth'))
```

```
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min,
cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=True,
pert_scale=1.0, fig_rows=2, fig_width=25, fig_height=11,
label_map=cifar_classes)
```

```
if device.type == 'cuda': torch.cuda.empty_cache()
```

```
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:557: U
serWarning: This DataLoader will create 4 worker processes in total. Our sugg
ested max number of worker in current system is 2, which is smaller than what
this DataLoader is going to create. Please be aware that excessive worker cre
ation might get DataLoader running slow or even freeze, lower the worker numb
er to avoid potential slowness/freeze if necessary.
```

```
warnings.warn(_create_warning_msg(
```



```
# Отразим отличия для fgsm_eps=(0.001, 0.02, 0.5, 0.9, 10) и выявим
закономерность/обнаружим отсутствие влияние параметра eps для сетей FC
LeNet на датасете MNIST, NiN LeNet на датасете CIFAR
fgsm_epss = [0.001, 0.02, 0.5, 0.9, 10]
```

```
for fgsm_eps in fgsm_epss:
```

```
    print(f"Используется fgsm_eps {fgsm_eps}")
    model = FC_500_150().to(device)
    model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
```

```
    display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min,
mnist_max, fgsm_eps, deep_args, has_labels=False, l2_norm=True,
pert_scale=1.0, fig_rows=2, fig_width=25, fig_height=11)
```

```
    if device.type == 'cuda': torch.cuda.empty_cache()
```

```
for fgsm_eps in fgsm_epss:
```

```
    print(f"Используется fgsm_eps {fgsm_eps}")
    model = FC_500_150().to(device)
    model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
```

```

    evaluate_attack(f'mnist_fc_fgsm_eps{fgsm_eps}.csv', 'results', device,
model, mnist_loader_test, mnist_min, mnist_max, fgsm_eps, is_fgsm=True)

    if device.type == 'cuda': torch.cuda.empty_cache()

for fgsm_eps in fgsm_epss:

    print(f"Используется fgsm_eps {fgsm_eps}")
    model = Net().to(device)
    model.load_state_dict(torch.load('weights/clean/cifar_nin.pth'))

    display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min,
cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=True,
pert_scale=1.0, fig_rows=2, fig_width=25, fig_height=11,
label_map=cifar_classes)

    if device.type == 'cuda': torch.cuda.empty_cache()

for fgsm_eps in fgsm_epss:

    print(f"Используется fgsm_eps {fgsm_eps}")
    model = Net().to(device)
    model.load_state_dict(torch.load('weights/clean/cifar_nin.pth'))

    evaluate_attack(f'cifar_nin_fgsm_eps{fgsm_eps}.csv', 'results',
device, model, cifar_loader_test, cifar_min, cifar_max, fgsm_eps,
is_fgsm=True)

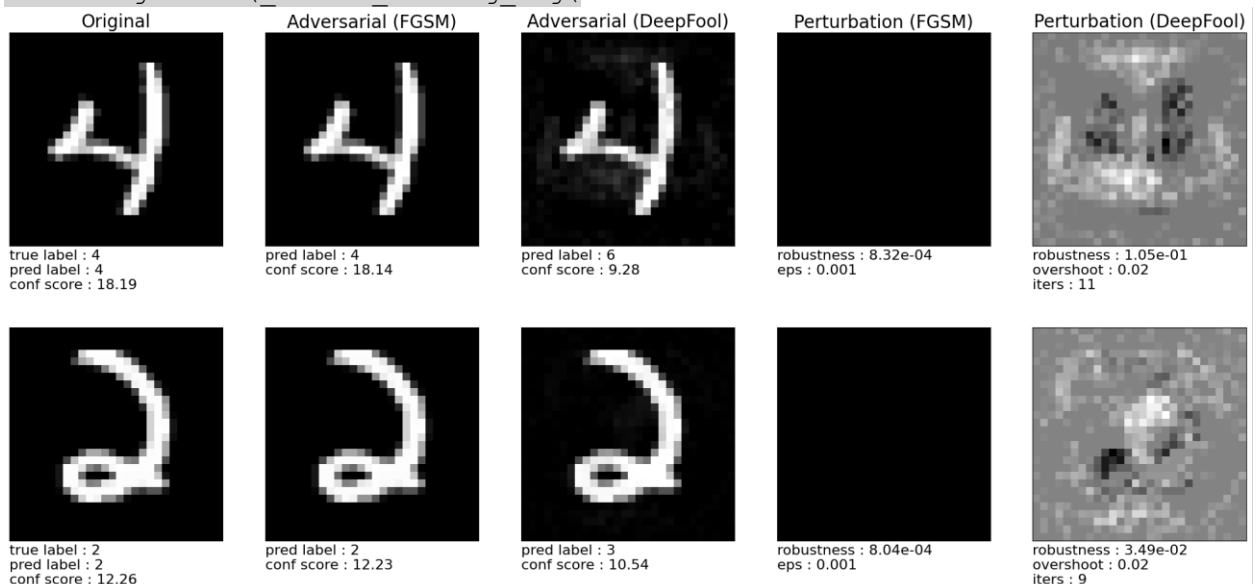
    if device.type == 'cuda': torch.cuda.empty_cache()

```


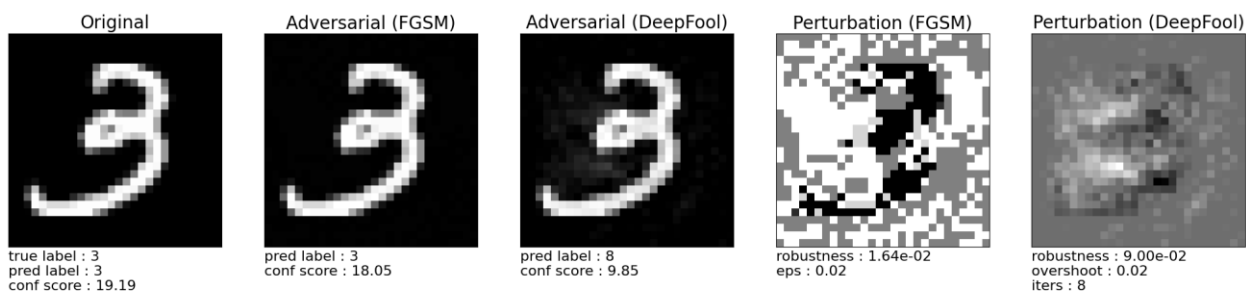
Используется fgsm_eps 0.001

/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:557: UserWarning: This DataLoader will create 4 worker processes in total. Our suggested max number of worker in current system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.

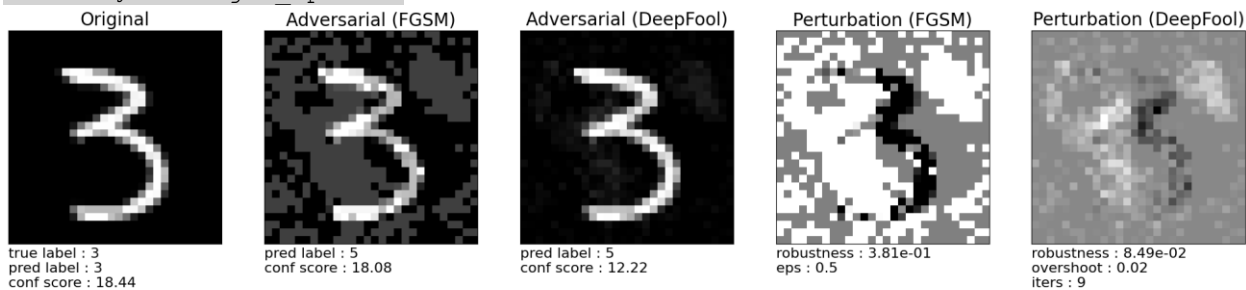
warnings.warn(_create_warning_msg(



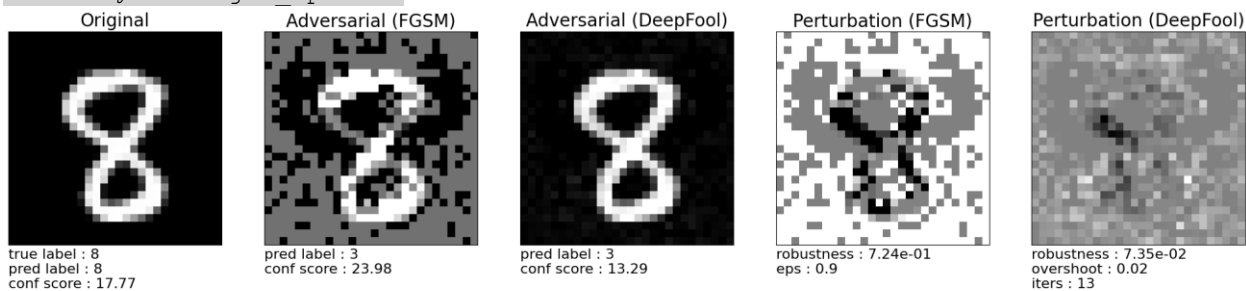
Используется fgsm_eps 0.02



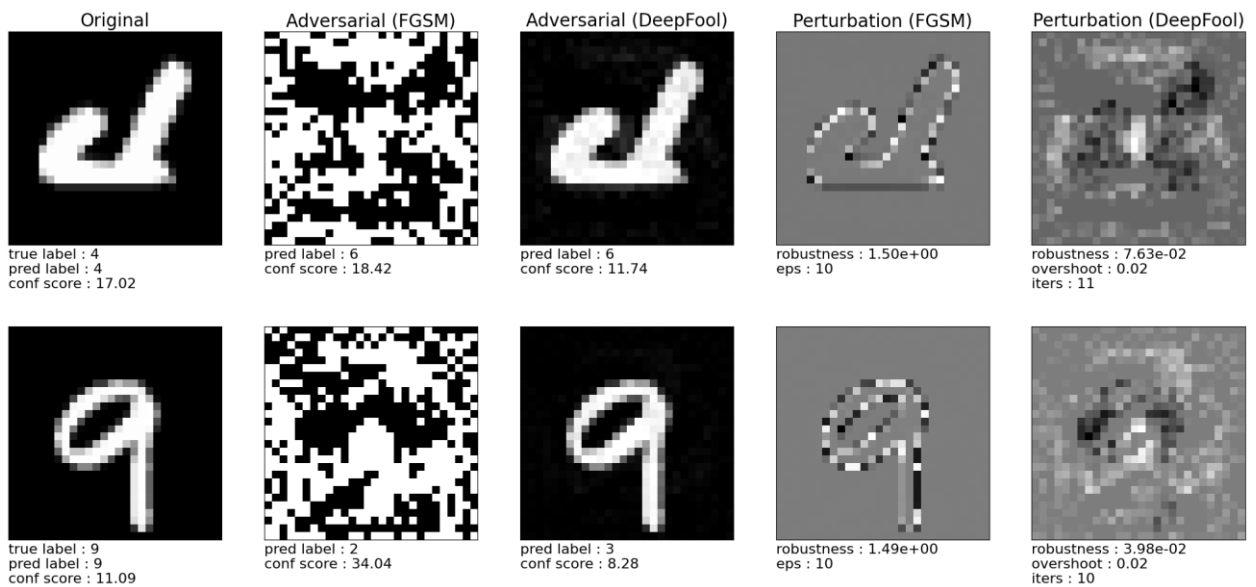
Используется fgsm_eps 0.5



Используется fgsm_eps 0.9



Используется fgsm_eps 10



Используется fgsm_eps 0.001

FGSM Test Error : 3.07%

FGSM Robustness : 8.08e-04

FGSM Time (All Images) : 0.72 s

FGSM Time (Per Image) : 72.00 us

Используется fgsm_eps 0.02

FGSM Test Error : 5.54%

FGSM Robustness : 1.60e-02

FGSM Time (All Images) : 0.54 s

FGSM Time (Per Image) : 53.68 us

Используется fgsm_eps 0.5

FGSM Test Error : 99.21%

FGSM Robustness : 3.86e-01

FGSM Time (All Images) : 0.56 s

FGSM Time (Per Image) : 56.40 us

Используется fgsm_eps 0.9

FGSM Test Error : 99.87%

FGSM Robustness : 6.86e-01

FGSM Time (All Images) : 0.60 s

FGSM Time (Per Image) : 60.29 us

Используется fgsm_eps 10

FGSM Test Error : 99.87%

FGSM Robustness : 1.47e+00

FGSM Time (All Images) : 0.55 s

FGSM Time (Per Image) : 55.18 us

Используется fgsm_eps 0.001

```
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:557: UserWarning: This DataLoader will create 4 worker processes in total. Our suggested max number of worker in current system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.
```

```
warnings.warn(_create_warning_msg(
```



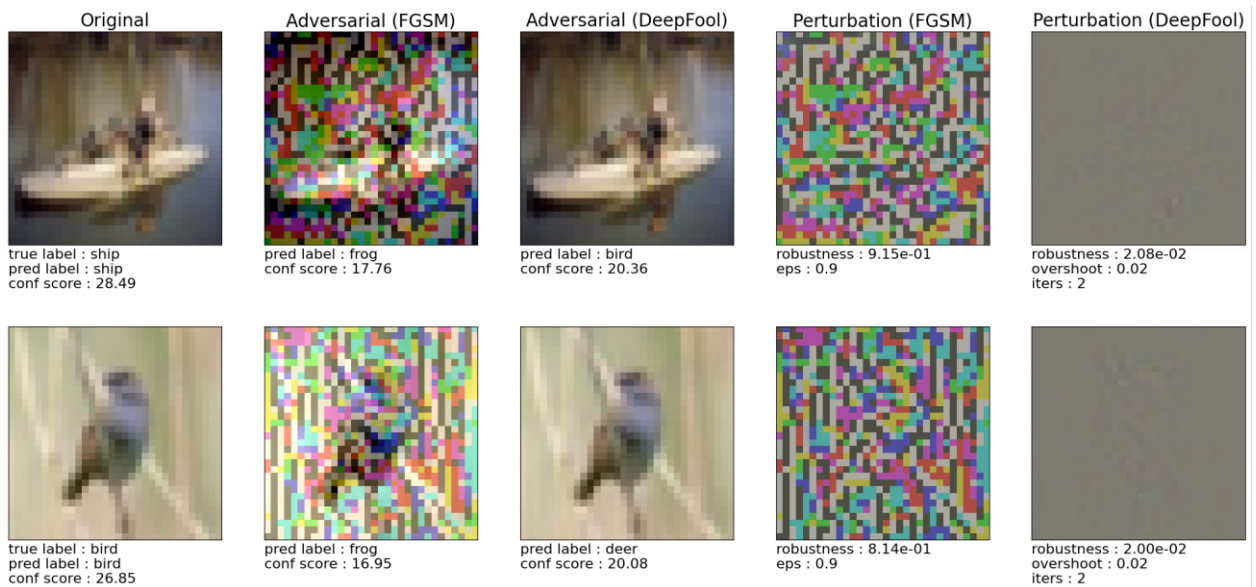
Используется fgsm_eps 0.02



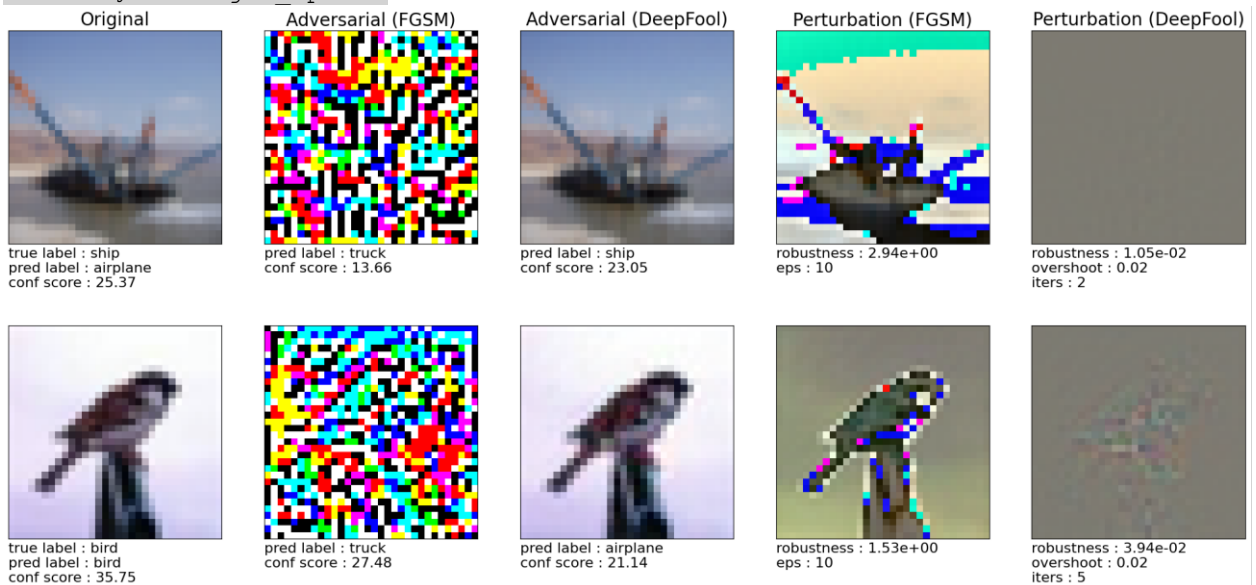
Используется fgsm_eps 0.5



Используется fgsm_eps 0.9



Используется fgsm_eps 10



Используется fgsm_eps 0.001

FGSM Test Error : 10.12%
FGSM Robustness : 8.92e-04
FGSM Time (All Images) : 1.09 s
FGSM Time (Per Image) : 108.82 us

Используется fgsm_eps 0.02

FGSM Test Error : 30.76%
FGSM Robustness : 1.78e-02
FGSM Time (All Images) : 1.31 s
FGSM Time (Per Image) : 131.25 us

Используется fgsm_eps 0.5

FGSM Test Error : 82.65%
FGSM Robustness : 4.40e-01
FGSM Time (All Images) : 1.12 s
FGSM Time (Per Image) : 112.16 us

Используется fgsm_eps 0.9

FGSM Test Error : 84.60%
FGSM Robustness : 7.79e-01
FGSM Time (All Images) : 1.16 s
FGSM Time (Per Image) : 116.42 us

Используется fgsm_eps 10

FGSM Test Error : 87.53%

```
FGSM Robustness : 2.46e+00
FGSM Time (All Images) : 1.15 s
FGSM Time (Per Image) : 115.36 us
```


```
import matplotlib.pyplot as plt

fgsm_eps = [0.001, 0.02, 0.5, 0.9, 10]
fgsm_test_error_MNIST = [3.07, 5.54, 99.21, 99.87, 99.87]
fgsm_robustness_MNIST = [8.08e-04, 1.60e-02, 3.86e-01, 6.86e-01, 1.47e+00]
fgsm_test_error_CIFAR = [10.12, 30.76, 82.67, 84.62, 87.50]
fgsm_robustness_CIFAR = [8.92e-04, 1.78e-02, 4.40e-01, 7.79e-01, 2.46e+00]

plt.plot(fgsm_eps, fgsm_test_error_MNIST, label='MNIST')
plt.plot(fgsm_eps, fgsm_test_error_CIFAR, label='CIFAR')

plt.xlabel('FGSM Epsilon')
plt.ylabel('Процент ошибок')
plt.title('Сравнение ошибки тестирования моделей')
plt.legend()

plt.show()
```

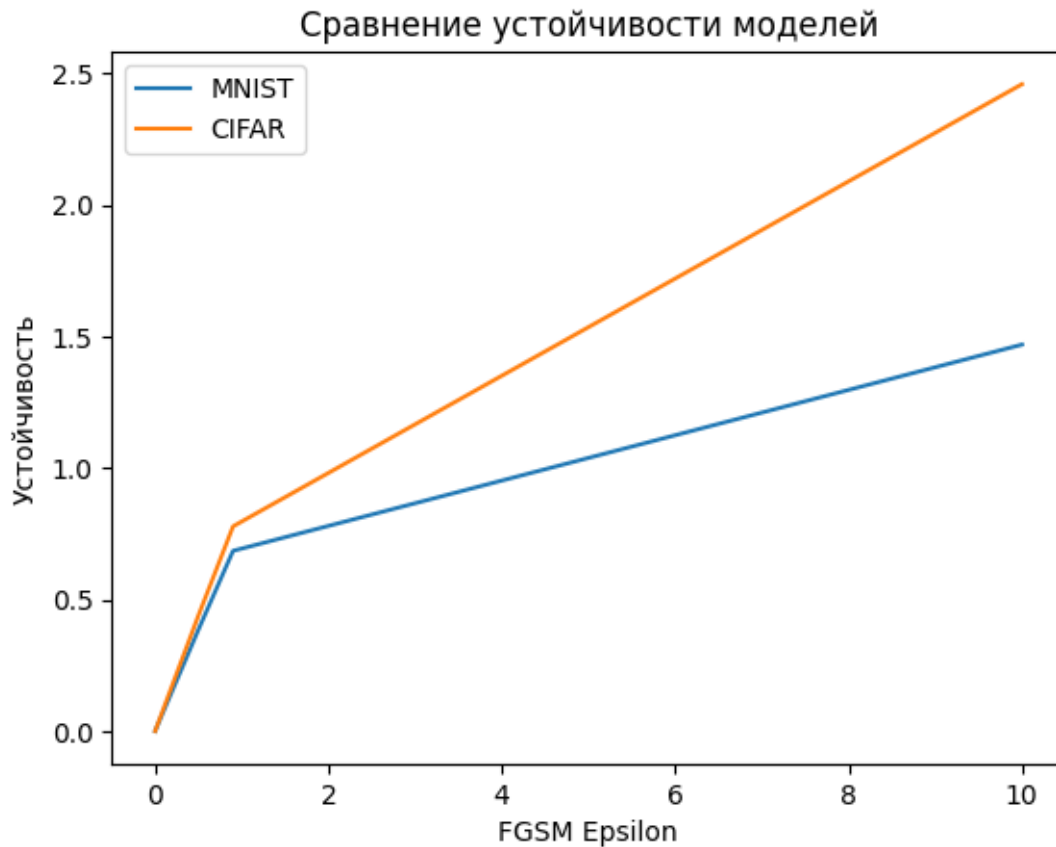


In [24]:

```
plt.plot(fgsm_eps, fgsm_robustness_MNIST, label='MNIST')
plt.plot(fgsm_eps, fgsm_robustness_CIFAR, label='CIFAR')

plt.xlabel('FGSM Epsilon')
plt.ylabel('Устойчивость')
plt.title('Сравнение устойчивости моделей')
plt.legend()

plt.show()
```



По картинкам понятно что от параметра `fgsm_eps` зависит степень шума. Параметр `fgsm_eps` имеет значительное влияние на производительность модели и ее устойчивость к атакам.

При увеличении значения `fgsm_eps` увеличивается ошибка тестирования модели, что указывает на снижение ее производительности.

В то же время, устойчивость модели также увеличивается с увеличением `fgsm_eps`, что указывает на то, что модель становится менее устойчивой к атакам. Это означает, что модель более подвержена ошибкам при обработке входных данных, которые были зашумлены.

Время вычислений не показывает последовательной тенденции с увеличением `fgsm_eps`. Это говорит о том, что величина возмущений не оказывает существенного влияния на время обработки изображений.